

1 Abstract

Paths on the Autonomous Systems (AS) graph of the Internet derived from the Border Gateway Protocol (BGP) can be used by researchers to understand the dynamics of Internet topology and to interpret how that topology may enable nations to enact censorship, surveillance, or other Internet control measures. Unfortunately datasets of these paths are not generally made publicly available, and paths collected from measurements such as traceroute measurements or BGP probes tend to be incomplete. Simulation frameworks have been used to generate AS paths based on empirical AS graphs. We have reimplemented one such framework, BGPSim [4] in Python as an efficient, open source BGP path simulation tool called BGPSimPy. With this tool we produced publicly available routing tree datasets that will enable future research on AS level topology. Previous research endeavors in this direction have only identified chokepoints in single snapshots. We apply the path simulation and routing tree datasets to view snapshots of the Internet over multiple years in order to introduce a new technique to investigate the evolution of the complex and dynamic AS level Internet topology. As an application of the routing trees we compare multiple qualitative sources of Internet freedom and censorship with chokepoint metrics in order to interpret the relationship between AS level topology and actual censorship activity.

2 Related Work

The primary struggle facing Internet topology researchers is a lack of ground truth data. Paths on the AS level of the Internet are incomplete and sometimes inaccurate [cite?](#). Additionally, AS business relationships are confidential and not publicly available. In this absence of solid empirical data, a technique for inferring AS relationships and a technique for inferring BGP paths are needed to study AS-level topology and evolution. The first effective relationship inference algorithm was developed by Gao in [3]. This algorithm was limited in that it focused on maximizing the occurrences of the valley-free rule on the AS graph, which has been shown to not always hold [cite?](#). Additionally, Gao's algorithm was verified against a single tier-1 provider's ground truth data. A more verifiable correct set of AS relationships was developed by the Center for Applied Internet Data Analysis (CAIDA) [6]. These relationships, which were generated using Gao's algorithm and other inference techniques, were verified against a larger ground truth dataset, and didn't require that the valley-free rule always hold. This leads to more realistic relationships. AS relationships do not tell the whole story of AS-level topology, however. In order to identify chokepoints in the AS graph, it is necessary to determine the paths that Internet traffic follows. Researchers either use a measurement tool like traceroute or BGP AS-path announcement data to determine AS paths. These techniques have been

shown to yield an incomplete and inaccurate set of paths [cite?](#). It is also important that paths from all ASes be utilized to get a complete picture of Internet topology. One technique for generating such a set of paths was developed by Gill et al. in [4]. The algorithm developed in [4], called BGPSim, uses a modified breadth-first search on an AS graph with defined relationships, as in from the CAIDA dataset. The breadth first search traverses the AS graph starting from a root node, and the search traverses edges according to priorities related to the relationships defined by Gao. The priorities, in the order they are integrated into the algorithm, are: Shortest Path (SP), Local Preference (LP), and Tie-break (TB). The resultant paths are ideal according to economic concerns of AS admins.

Previous studies have used BGP path models to find ASes that intercept a high fraction of paths. One such project in [2] identified that 90% of paths on the Internet could be intercepted with only 30 or so ASes. The researchers in [2] generated paths by first using only paths to top websites as defined by the Alexa top websites project, and then appending additional edges to those paths from the full AS graph according to the principles defined by Gao [3]. Many of the ASes that were found to intercept a large number of paths were found to be within nations that conduct censorship. As an extension of these results, another paper [5] revealed that ASes that intercept many paths could be utilized for decoy routing. This approach for identifying AS level chokepoints has several potential pitfalls that are remedied by the approach taken in this paper. First, the authors in [2] didn't consider that many of the paths found from the Alexa top websites would have destinations in nations that censor Internet traffic, particularly China due to its large Internet population. This artificially inflates the chokepoint nature of Chinese ASes. As an alternative, in this paper we consider paths from every source-destination AS pair, and we make a distinction between in-to-out paths and out-to-in paths. Secondly, chokepoints have previously only been identified at a single snapshot of the Internet. This makes it difficult to discuss the evolution of Internet topology, and whether or not chokepoints are anomalous or common. Finally, the aforementioned approach for identifying chokepoints does not allow a clear comparison between different nations. It can be said that one nation controls a large portion of Internet paths, but not how easily traffic directed through that country could be intercepted on a national level. For this, some aggregate measure across all of the border ASes within a country must be considered, as it is in this paper.

In [7], Xu et al. investigated the AS level topology of China to identify where keyword filtering occurred. They found that the most effective ASes with which to deploy keyword filtering devices are those in the backbone of the Chinese AS topology. A relevant contribution of [7] is that, while most filtering occurs in border ASes, some filtering is controlled by non-central provincial ASes. China has a diverse strategy for Internet censorship, both targeting chokepoints and

the Chinese provincial network. The potential for various forms of censorship in regards to various AS level topologies motivates the question: Is centralized censorship or decentralized censorship more common? Instead of directly identifying censorship devices on the AS graph, we instead quantify the chokepoint potential of ASes on the national level, and then compare that with qualitative Internet freedom measures and empirical censorship events.

3 BGPSimPy

In order to conduct an analysis on AS level topology, the paths generated by the Border Gateway Protocol (BGP) have to be modeled. A traditional approach is to use a tool like traceroute to collect paths via measurement, or at a larger scale use a selection of BGP data collectors with a tool like BGPStream [?] [?]. Unfortunately, these techniques require extensive collection time and are still susceptible to errors and incomplete data. There is no readily available BGP path dataset with which to study. Thus, researchers must turn to simulation to find accurate BGP paths.

One such simulation framework, the previously mentioned BGPSim, was developed by Gill et al. [4] as a tool for quickly estimating AS paths from an input AS graph. This AS graph could be taken from a set of inferred AS relationships, such as the CAIDA AS relationship dataset [1]. BGPSim outputs routing trees, which contain all paths from leaf ASes to the root AS, for each AS in the AS graph. The routes are selected using a modified breadth-first search that selects with the following priorities: shortest path, local preference (using the rules from Gao [?]), and finally tie break of all equally good paths. With these priorities, BGPSim generates a realistic collection of AS level paths, which are ideal for use in Internet topology studies.

Unfortunately, BGPSim is not ready to deploy immediately and easily, as would be necessary for a tool dedicated to allow fast monitoring of the Internet's overall state. This is largely due to the fact that BGPSim is written in C# (affecting its ability to be used in a cross-platform way) with DryadLinq (a deprecated C# library) for distributed computing. As a remedy to these problems, we have reimplemented BGPSim under the name BGPSimPy. BGPSimPy is written purely in python, making it cross platform and easy to integrate. Additionally, distributed computing is now achieved with MPI using mpi4py, a framework that is supported on many computer clusters.

4 Generating Routing Trees

AS level Internet topology research is plagued by a lack of publicly available path data. Path inference is also computationally expensive for the entire Internet, prohibiting continuous research and Internet monitoring. Thus, we have chosen to release the routing tree datasets generated for this paper. This is a major contribution to Internet research, and will allow new en-

deavors to investigate Internet topology relationships on an unprecedented scale. As a focus of this paper, and an example application, we use these routing trees to examine topological evolution of the Internet and its relationships to Internet freedom and censorship.

5 Applications of Routing Trees

5.1 Measuring Chokepoints in AS-level Paths

Once we generated routing trees, we could use the paths to identify chokepoints, or ASes that intercept large portions of paths, on the AS graph. Specifically, we chose to evaluate border ASes, or ASes that could reach an AS registered to a different nation with one hop. By selecting border ASes we could locate the ASes that intercepted traffic into a nation (via out-to-in paths), as in a foreign attempt to access a domestic website, or ASes that intercepted traffic out of a nation, as in a domestic attempt to access a foreign website (via in-to-out paths). We choose to make a distinction between these two path types to better understand the implications of the chokepoints belonging to each nation.

To evaluate border ASes we use the metric of chokepoint potential, defined here to be the ratio of the number of paths intercepted by a border AS to the number of paths for that nation, as shown in equation ???. This metric is calculated separately for both in-to-out paths and out-to-in paths. This calculation is performed by traversing the routing trees and keeping record of which border nodes intercept each path, as well as what path type is currently being traversed with regards to the host nation of the border AS in question. Once all paths are traversed, the totals are compiled together, with a chokepoint potential value stored for each border AS.

The chokepoint potentials of the border ASes allowed us to investigate the state of the AS-level topology of the Internet for any given set of initial AS relationships. We first looked at how many ASes each nation is required to control to intercept a percentage of paths. For this test, we chose 90%, as this would indicate a clear control of Internet paths for the given country and was also the percentage used in [2]. Additionally, we determined this value for each nation for multiple timestamps, yielding our initial picture of the evolution of Internet topology. The results of this test for some nations of interest are shown in figure 1.




Figure 1: Number of border nodes needed to control 90% of paths for various nations

These results give an idea of how the Internet has evolved for different nations. Some nations, particularly China, have evolved to a state where it is significantly easier to intercept paths. Others, for instance Brazil, have become more open with regards to the chokepoint nature of border ASes. These results are even more striking when one considers that the number of ASes globally has drastically increased over the years, amplifying the effect of higher chokepoint potentials. The number of ASes in the CAIDA AS relationship dataset in 2012 was 43,274 in January, while in January of 2018 the count was 60,006. Next, we took a more fine-grained approach to investigate the dynamics and evolution of AS-level chokepoints. For each nation, we sorted the border ASes in descending order into a list, so that the first AS in the list was the AS with the highest chokepoint potential. Then, we selected each AS from the list and kept track of the ratio of paths intercepted. This test allowed us to investigate the ease with which different nations could control paths at multiple timestamps. Additionally, this experiment illuminates that the topological state of a nation in regards to border ASes is often complex and dynamic. Since nations possess wide variations in AS count, this data is best visualized on a semi-logarithmic scale. The different timestamps are shown in different colors, where red is the furthest in the past and blue is the most recent. These results are depicted in 2.

The variation between the different nations highlights the importance of measuring chokepoints on the national level. Trends exist within a nation, but nations tend to have surprisingly unique AS graph dynamics. China shows a striking increase in chokepoint potential. In order to control 90% to utilize a strikingly few number of ASes. This suggests an ever strengthening backbone in the Chinese Internet. Conversely, the trend that Brazil’s capability to intercept paths is decreasing is also evident in these results. This could be the effect of an expansion of infrastructure, leading to more border ASes and more diverse paths in and out of the country. [Maybe there is a source to support this?](#) Another interesting element of these results is

the dynamic behavior for some nations. It seems that in some cases, when it becomes easier to control many paths, it becomes more difficult to control all paths. [Not sure of implications, needs work here](#)

5.2 Comparing Internet Topology with Freedom On The Net

While research in the past [2] has connected AS-level chokepoints with the capability of a nation to conduct censorship, linking the topology based measurements to other attempts at measuring Internet freedom has not been done before to the best of our knowledge. Such a connection would bridge the gap between a dynamic system driven by many factors, both automatic and instituted, and qualitative classification of nations.

Freedom House’s Freedom On The Net (FOTN) report [?], is an annual ranking and analysis of Internet freedom, interpreted from multiple measures including obstacles to access, limits on content, and violations of user rights. The FOTN includes a ranking of nations by a freedom index, where a higher number indicates a less free Internet for that nation. The index is calculated from a series of questions that are evaluated according to various sources regarding each nation. These questions encompass multiple categories, including government filtering and censorship of web traffic. In order to investigate whether the topology of a nation’s ASes is related to Internet freedom, we compare the FOTN freedom index with the ease with which a nation can control AS level paths. In figure 3 we plot the number of ASes required for a nation to control 90% of paths with that nation’s freedom index for multiple timestamps of AS relationships.

For each date we fit the data with an ordinary least squares linear regression, and found the relationship to be statistically significant (all p-values are less than 0.05). [need some better way to show exact p-values](#) The relationship shows that more free nations, as determined by FOTN scores, tend to have a more difficult to control AS-level topology, as in border nodes that control fewer portions of paths in and out of that country. This suggests that either the desire of national governments to control Internet traffic has impacted the layout of AS relationships, the topology of the modern Internet has evolved in a way to enable more successful filtering and blocking, or a combination of both. Up to this point it has been assumed that Internet topology and Internet censorship are related, but this has not been supported by both a measure of Internet freedom and actual AS- level paths.

5.3 Comparing Internet Topology with Empirical Censorship Data

Another technique for relating our chokepoint metric with censorship is to look at empirical censorship data. One such dataset is provided by the Open Observatory of Network Interference (OONI) [?]. OONI is a project wherein Internet users around the world use a tool (the ooniprobe) to examine evidence of Internet censorship.

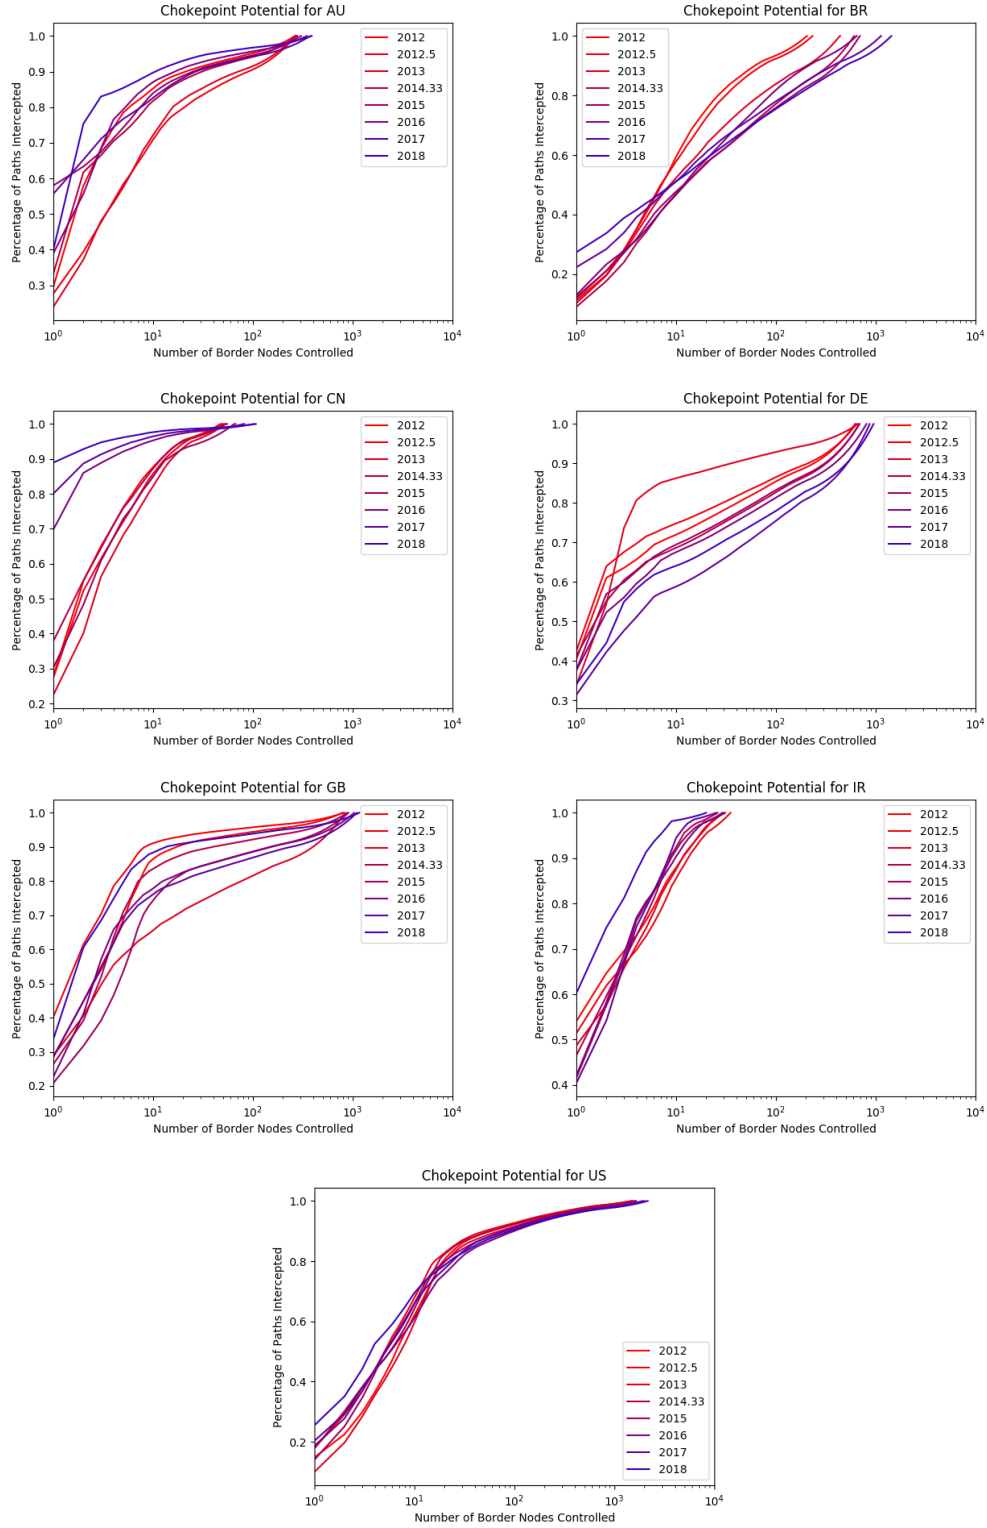


Figure 2: Ratio of Paths Intercepted for a certain number of Border Nodes for multiple nations

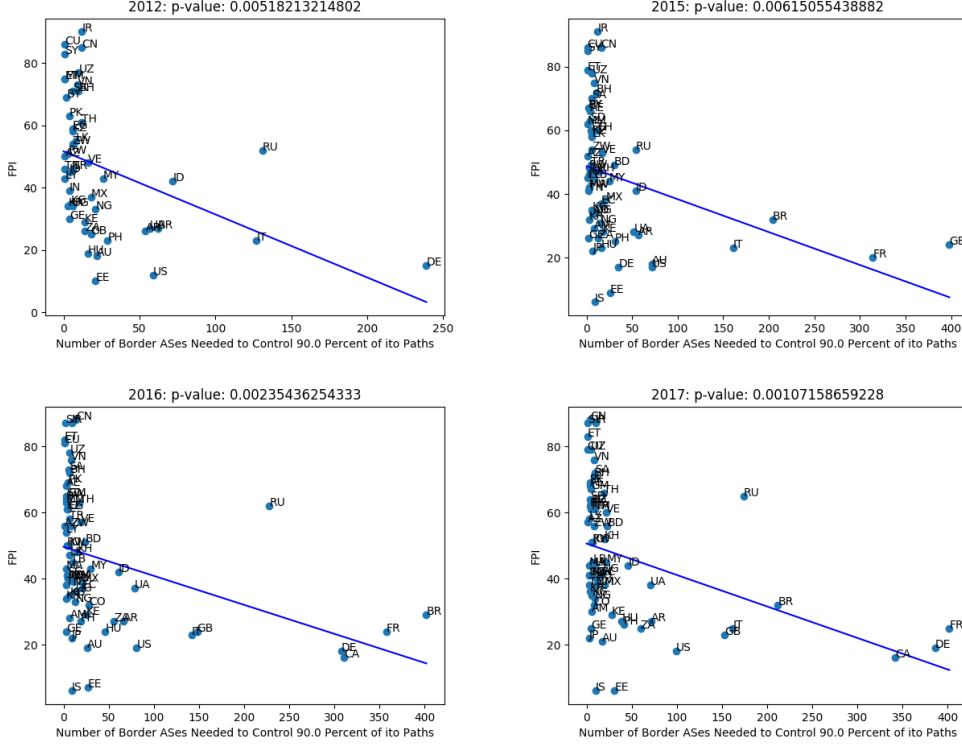


Figure 3: Ratio of Paths Intercepted for a certain number of Border Nodes for multiple nations

The ooniprobe allows users to run a series of tests, such as an HTTP requests that might be known to be blocked by certain nations. The ooniprobe then compares the results of tests with previously determined unblocked examples of the same tests. Results where the test differs from expected are considered anomalies, and further heuristics can determine if the anomalies are likely to be censorship events.

We examined 18 days of OONI data from January in 2018. We recorded the ratio of anomalies in measurements for several classes of anomalous behavior to total measurements carried out in each nation. We used the web connectivity test from OONI that records the following fields to measure connectivity: DNS, HTTP_FAIL, HTTP_DIFF, and TCP. Each of these fields are recorded as true or false, according to whether the test resulted in a potential censorship event. We then plotted the ratio of potential censorship events in the dataset against a national metric of chokepoint potential. In this case, we again looked at the number of border nodes required to intercept 90% of AS-level paths. The results of this test are plotted in figure 4. A relationship like the one seen in the FOTN data is not as clear, perhaps due to the uneven sampling across nations. It is clear, however, that nations can be classified both by its chokepoint potential and the occurrence of possible censorship events. A na-

tion, for instance, with a high ratio of web connectivity anomalies with a high chokepoint potential can be considered a censorship active nation with a strongly hierarchical AS topology. Thus, this analysis serves as a monitoring tool for the state of nations on the Internet in regards to actual censorship implementation, and potential for further censorship.

6 Discussion

We have introduced an efficient technique for studying the AS-level topology of the Internet. Additionally, we introduced the metric of chokepoint potential in order to evaluate the control that border nodes possess over paths into and out of a nation. We have applied the routing trees generated by BGPsimPy along with the chokepoint potential metric to present an overview of the evolution of the amount and strength of chokepoints on the AS graph over the past 6 years. We also connected our analysis of AS-level chokepoints in a novel way to both qualitative measures of Internet freedom and empirical censorship data. The former comparison reveals a strong relationship between the chokepoint potential of a nation and its lack of Internet freedom, while the former provides insight into the categories of nations, in regards to their implementations of censorship and the layout of their AS relationships.

References

- [1] Caida as relationships. <http://data.caida.org/datasets/as-relationships/>.

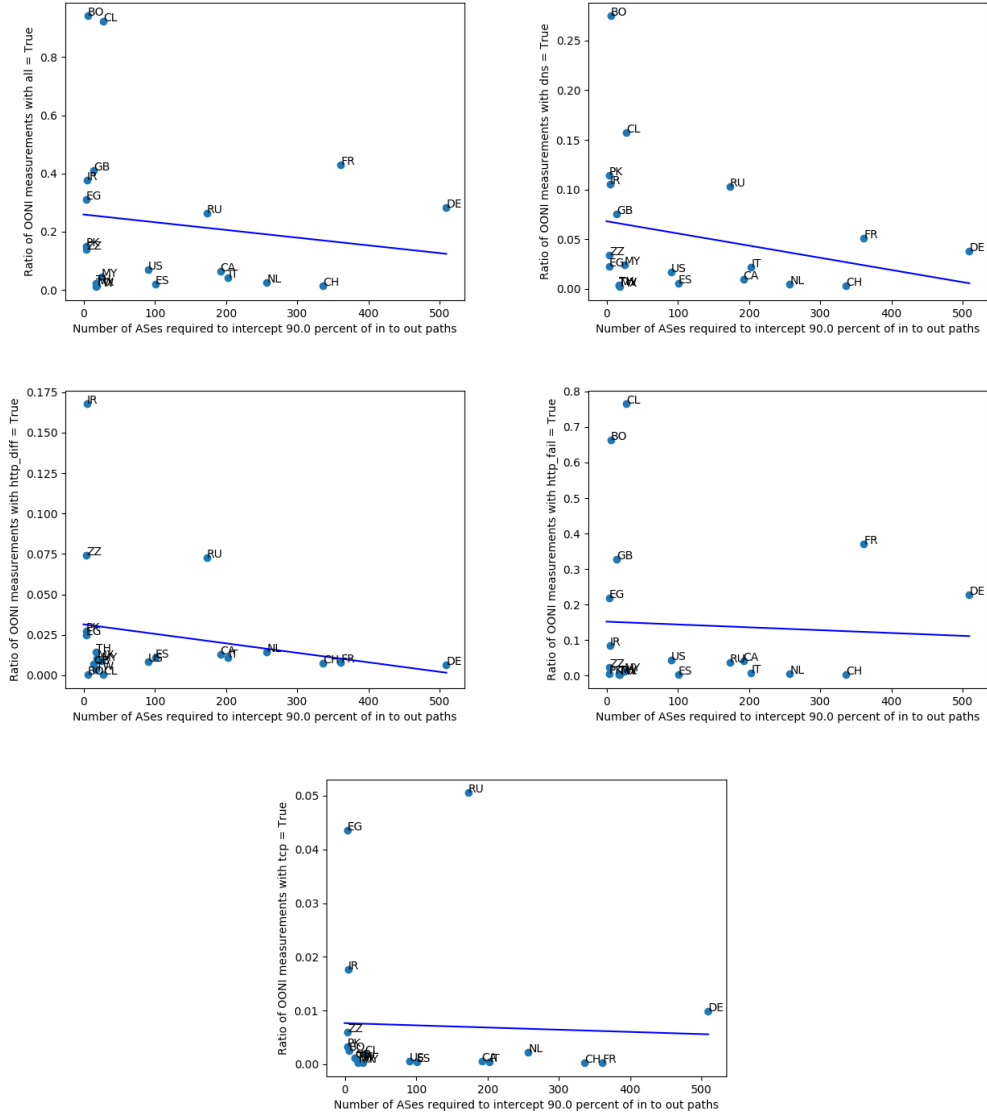


Figure 4: Ratio of anomalies in OONI measurements compared to chokepoint potential

- [2] HB Acharya, Sambuddho Chakravarty, and Devashish Gosain. Few throats to choke: On the current structure of the internet. In *Local Computer Networks (LCN), 2017 IEEE 42nd Conference on*, pages 339–346. IEEE, 2017.
- [3] Lixin Gao. On inferring autonomous system relationships in the internet. *IEEE/ACM Transactions on networking*, 9(6):733–745, 2001.
- [4] Phillipa Gill, Michael Schapira, and Sharon Goldberg. Modeling on quicksand: dealing with the scarcity of ground truth in interdomain routing data. *ACM SIGCOMM Computer Communication Review*, 42(1):40–46, 2012.
- [5] Devashish Gosain, Anshika Agarwal, Sambuddho Chakravarty, and HB Acharya. The devil’s in the details: Placing decoy routers in the internet. In *Proceedings of the 33rd Annual Computer Security Applications Conference*, pages 577–589. ACM, 2017.
- [6] Matthew Luckie, Bradley Huffaker, Amogh Dhamdhere, Vasileios Giotsas, et al. As relationships, customer cones, and validation. In *Proceedings of the 2013 conference on Internet measurement conference*, pages 243–256. ACM, 2013.
- [7] Xueyang Xu, Z Morley Mao, and J Alex Halderman. Internet censorship in china: Where does the filtering occur? In *International Conference on Passive and Active Network Measurement*, pages 133–142. Springer, 2011.